

Reverse Engineering

Josh Hofing

Last Week:

- Basics of x86_64
 - Registers
 - Memory
 - Instructions
 - The Stack
 - Calling Conventions

This Week:

- Virtual Memory
- Page Tables
- The OS
- Theorem Provers

Virtual Memory

```
main:
00400895 55          push    rbp
00400896 4889e5     mov     rbp, rsp
00400899 4883ec10   sub     rsp, 0x10
0040089d b800000000 mov     eax, 0x0
004008a2 e866ffffff call    init
004008a7 bfa80a4000 mov     edi, 0x400aa8 {"Can you tell me where to mail th..."}
004008ac e84ffdffff call    puts
004008b1 b800000000 mov     eax, 0x0
004008b6 e876ffffff call    get_number
004008bb 488945f0   mov     qword [rbp-0x10], rax
004008bf 488b45f0   mov     rax, qword [rbp-0x10]
004008c3 488945f8   mov     qword [rbp-0x8], rax
004008c7 488b45f8   mov     rax, qword [rbp-0x8]
004008cb 488b10     mov     rdx, qword [rax]
004008ce 48b8eecefa0d000d... mov     rax, 0xd00dfaceee
004008d8 4839c2     cmp     rdx, rax
004008db 751b     jne     0x4008f8
```

```
004008f8 bf000b4000 mov     edi, 0x400b00 {"That doesn't look right... try a..."}
004008fd e8fefcffff call    puts
00400902 b801000000 mov     eax, 0x1
```

```
004008dd bfd80a4000 mov     edi, 0x400ad8 {"Got it! That's the right number!"}
004008e2 e819fdffff call    puts
004008e7 b800000000 mov     eax, 0x0
004008ec e88cfeffff call    print_flag
004008f1 b800000000 mov     eax, 0x0
004008f6 eb0f     jmp     0x400907
```

```
00400907 c9          leave   {__saved_rbp}
00400908 c3          retn
```

Virtual Memory

- The address here is interesting...
- Is the first instruction of main *really* at 0x00400895?
- What if I run 2 copies of the program?
- What if I run 2 different programs?

```
main:
00400895 55                push    rbp
```

Virtual Memory

- Most modern hardware supports *Virtual Memory*
- The addresses we tell the program it is running at are different than the physical locations in memory they live at
- How do we tell the hardware how a vaddr maps to a physical address?

Page Tables

- Mapping of a *Page* of Virtual Addresses to a *Page* of Physical Addresses
- How big is a Page?
 - It depends on the architecture
 - We'll only really ever look at 4KB pages

Page Tables

- Pages are *aligned*, meaning the bottom bits are all zeroes
- So, to map the region containing 0x00400895, you map 0x00400000 to 0x00400FFF
 - $0xFFF = 4KB - 1$

Page Permissions

- Pages have permissions associated with them:
 - Read (R) = Can I read data from the page?
 - Write (W) = Can I write data to the page?
 - Execute (X) = Can I execute instructions on the page?
- This will matter when we talk about pwnning!

The OS

- Modern Operating Systems are really complicated!
- They deal with input and output to a huge number of devices
 - Keyboards/Mice, Networking, GPUs, etc.
- They manage a bunch of programs running at once
- They manage page tables

Talking to the OS

- When a program wants the OS to do something, it issues a *System Call*
 - The *syscall* instruction does this on x86_64
- Registers are setup, just like function arguments
 - rax contains the *Syscall Number*
 - rdi, rsi, rdx, r10, r8, r9 are arguments

Syscalls

- There are lots of syscalls on modern Linux (300+)
 - And C functions that mimic the common ones!
- The main uses are things like dealing with files, networking, and allocating memory

Files (and things like them)

- In Unix-like operating systems, everything is a file
 - Files (duh)
 - Sockets (for networking)
 - Information about processes (/proc)
 - Attached devices (/dev)
 - System settings (/sys)

File Descriptors

- When you use the *open* syscall, you are given back a *handle* to a file, also known as a *file descriptor*
- A file descriptor is just a number
 - Some file descriptors that come built-in to the program:
 - 0 = standard input (stdin)
 - 1 = standard output (stdout)
 - 2 = standard error (stderr)

The network is files, too!

- The *socket* syscall creates a new socket
 - You can use it as a client or a server

Using Files

- There are 3 main operations all files support:
 - *read, write, and close*
- Same API for files, network, system information, randomness, ...
- Pretty neat!

Example

- <https://godbolt.org/g/7ijzxF>
- this will help: <https://filippo.io/linux-syscall-table/>

Theorem Proving

- Sometimes, programs involve complicated instructions
- *recurse* is a great example

Z3

- We imported from z3, which is a fantastic *theorem prover*
- Give it a set of constraints, and it will produce values that satisfy them (or say they are *unsatisfiable*)
 - Sometimes this will take awhile...
- A commonly used tool in CTFs

Ints, BitVecs, Solvers

- Z3 supports a bunch of types that it understands
 - You'll mostly see:
 - Ints (arbitrary-size integers)
 - BitVecs (Integers of a specific bit-length)
 - Bools (True or False)
 - a Solver (the API for checking constraints)

Basic Example

```
from z3 import Ints, Solver
a, b = Ints('a b')
s = Solver()
s.add(a + b == 1234)
s.add(a - b == 500)
print(s.check())
print(s.model())
```

```
>>> sat
```

```
>>> [a = 867, b = 367]
```

Another Example

```
from z3 import BitVecs, Solver
a, b = BitVecs('a b', 16)
s = Solver()
s.add(a ^ b == 0xbeef)
s.add(a == 0xdead)
print(s.check())
print(s.model())
```

```
>>> sat
```

```
>>> [b = 24642, a = 57005]
```

Recurse

```
int recurse(int a, int b, int c) {  
    int sum = a + b;  
    if (c == 16 && sum == 116369) {  
        return 1;  
    } else if (c < 16) {  
        return recurse(b, sum, c + 1);  
    } else {  
        return 0;  
    }  
}
```

And in Python

```
def recurse(a, b):  
    for _ in range(17):  
        a, b = b, a + b  
    return b
```

```
a, b = raw_input().split(' ')  
a, b = int(a), int(b)  
assert(recurse(a, b) == 116369)
```


Solving without doing math!

```
from z3 import Ints, Solver
a, b = Ints('a b')
rec = recurse(a, b)
s = Solver()
s.add(rec == 116369)
print(s.check())
print(s.model())
```

```
>>> sat
```

```
>>> [b = 37, a = 13]
```

This is really powerful

- Doing the solve is also super NP-Complete
- If your equations get too big or too complicated, you're gonna have a bad time
- I've had equations that Z3 spent hours trying to solve
- It's a powerful tool, but not always the most scalable