

CS-3943-G / CS-9223-H

## Introduction to ROP

# Pwning so far

- So far all the binaries we've pwned have either:
  - Had NX off, allowing us to write shellcode
  - Had a “give\_shell” function
  - Had “system” imported

# Pwning so far

- Most real programs don't have any of these properties
  - So they're secure, right? We're all out of a job!
  - Time to retire...
  - Just kidding, we can still win against the bugs!

# Return Oriented Programming

- ROP (Return Oriented Programming) to the rescue!
  - ROP is a Code-Reuse attack, which means we don't write our own code to insert into the program
  - Instead, we use the bytes already in the program, which are marked executable
  - Of course, the clever programmers have gotten rid of all the useful functions, like “system”, so we'll need to be clever

# Some observations about x86

- Functions aren't really a “thing” in x86...
  - You can jump to the middle of a function, and that's all fine!
- Instructions are variable-length in x86...
  - You can jump into the middle of an instruction, and that's fine!
  - Jumping at different points yields different opcodes, which means we have a lot more instructions to use than the programmer intended
- If we can smash the callstack, we can control return pointers
  - That means we can do evil things like return, then return again, etc...

# Functions in x86

- x86 *kinda* understands functions
- What it really understands is the call stack
  - *call* and *ret* instructions push and pop to the callstack, for instance
- What does the *ret* instruction really do?
  - *pop rip*
  - If we control the stack, we control what it pops
    - That means we can control multiple returns, each one making small changes

# Instructions in x86

- Instructions are multiple bytes long, and they're variable length
  - Many instructions, like *ret*, are 1 byte (opcode = c3)
  - That means, if you look for all instances of c3, you can find sequences of instructions that lead up to a *ret*
- You can just jump to any byte, and the processor will happily decode it and run it...
  - So you can jump into the *middle* of multi-byte instructions

# Example

- `mov rax, 0xc30f05 ; ret` == 48 c7 c0 05 0f c3 00 c3
  - If we offset a few bytes into this instruction sequence...
    - 05 0f c3 == “`syscall ; ret`”
- That’s a pretty powerful example. Maybe a little contrived...
  - In a large enough program you’ll have tons of these little “gadgets” to do fun things with
- There are programs that find these gadgets for you:
  - I like `rp++`, some people swear by `ROPGadget`



# pop ; pop ; pop ; ret

- Remember, arguments are passed in registers!
  - If we want to pass arguments to functions, we need to set registers
- Luckily, all the *pop <register>* instructions are 1-byte long
  - Normal programs might even have them leading up to a *ret*
- If we want to set the first argument, we might find a *pop rdi ; ret* gadget, and then have our next return go to the function.
  - [ pop\_rdi\_ret, hello\_world, puts ]

# Figuring out where libc is

- On modern systems, we put shared libraries in a different place every time the program starts
  - That way, we can't just predict where "system" will be loaded and jump to it
- Some part of our program needs to know where it is, to call functions...
  - The GOT holds pointers to libc once they get resolved!
  - If we can read some bytes out of the GOT... we can know where libc is!
  - And since libc offsets from the base are always the same...
    - We can calculate any address in libc from a single GOT leak!

# Imagine this exploit...

- We have full control of the stack
  - The program did `gets(buf)`, or something similarly vulnerable
- We start our ROP...
  - `puts(puts_got_address)` to leak it
  - This prints out the address of `puts` in `libc`
  - We subtract the offset of `puts` in `libc` from that address to get the base address of `libc`
  - Then we add the offset of `system` in `libc` to get the real address of `system`
  - We jump there, with an argument of `"/bin/sh"`, and get a shell
  - Cool!