

# ROP

CS-UY 3943-G / CS-GY 9223-H

# Motivation

- Consider a situation where we have:
  - A buffer overflow or other vulnerability that gives us control over the return address and values on the stack
- But:
  - No executable stack (NX)
  - No handy shell-like functions (give\_shell, run\_cmd, system)
- How can we still get code execution?

# Return-Oriented Programming

- Solution: use code that's already in the program!
- x86 code is *variable length* (each instruction can be 1 to 15 bytes long) and *dense* (most random byte sequences are valid instructions)
- This means that you can jump to the middle of an instruction and usually get a valid instruction

# Pop, pop, ret

- The core idea is to find sequences of instructions (called *gadgets*) in the binary that end in *ret*
- Each of these does a little bit of work and then transfers control to the next gadget via *ret*

# Simple Example: `execve("/bin/sh")`

- Let's give ourselves some gadgets, and assume the string `"/bin/sh"` is in memory at `0x600000`
- The full call to `execve` is `execve("/bin/sh", NULL, NULL)`
- In assembly, we roughly want

```
mov rdi, 0x600000
mov rsi, 0
mov rdx, 0
mov rax, 59
syscall
```

# Simple Example: `execve("/bin/sh")`


- Our gadgets:

A	B	C	D
syscall	pop rdi	pop rsi	pop rdx
ret	ret	ret	ret

E
pop rax
ret

- What should our stack look like?



# Simple Example: `execve("/bin/sh")`

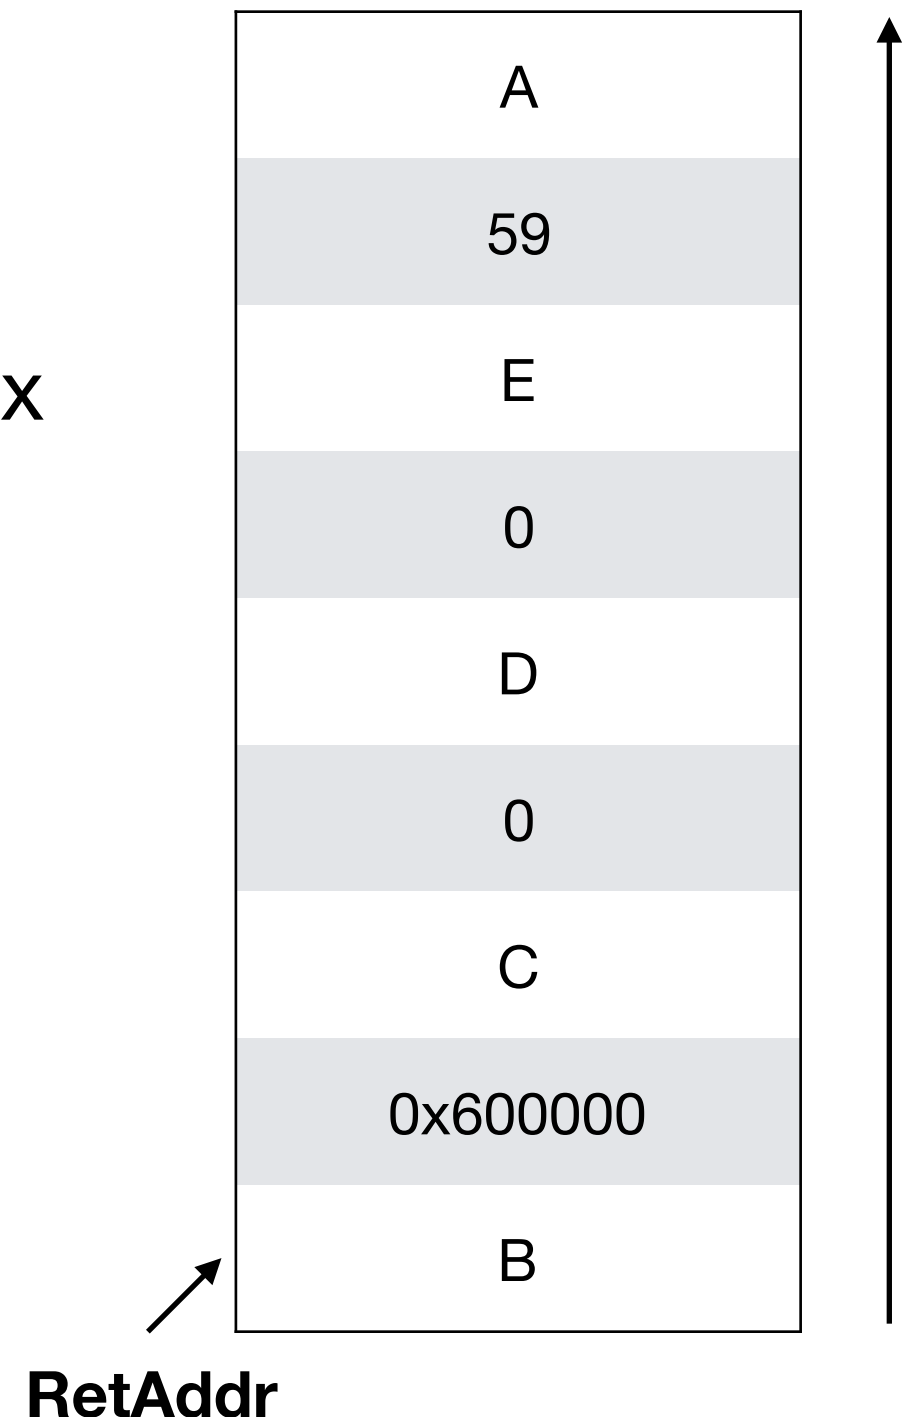
- Our gadgets:

A	B	C	D
syscall	pop rdi	pop rsi	pop rdx
ret	ret	ret	ret

E
pop rax
ret

- What should our stack look like?



# How do we find gadgets?

- Reminder: there are utilities that can do this for you!
- One that comes with pwntools is ROPgadget

```
(pwn) moyix@lorenzo:~/offsec/week_9/rop$ ROPgadget --binary ./rop
```

```
Gadgets information
```

```
=====
0x0000000000400564 : adc byte ptr [rax], ah ; jmp rax
0x00000000004005a4 : adc byte ptr [rax], ah ; jmp rdx
0x000000000040062d : add al, ch ; retf -1
0x00000000004006bf : add bl, dh ; ret
0x000000000040062b : add byte ptr [rax], al ; add al, ch ; retf -1
0x00000000004006bd : add byte ptr [rax], al ; add bl, dh ; ret
0x00000000004006bb : add byte ptr [rax], al ; add byte ptr [rax], al ; add bl, dh ; ret
0x00000000004006bc : add byte ptr [rax], al ; add byte ptr [rax], al ; ret
0x00000000004004a3 : add byte ptr [rax], al ; add rsp, 8 ; ret
0x00000000004006be : add byte ptr [rax], al ; ret
0x00000000004005c8 : add byte ptr [rcx], al ; ret
0x00000000004005c4 : add eax, 0x200a8e ; add ebx, esi ; ret
0x000000000040058b : add eax, edx ; sar rax, 1 ; jne 0x40059c ; pop rbp ; ret
0x00000000004005c9 : add ebx, esi ; ret
[...]
```




# Dealing With LibC

- Many binaries don't have a lot of gadgets available
- You may want to call interesting functions in libc that aren't imported by the target binary (so they're not in the PLT/GOT)
- Unfortunately, the address of libc is *randomized*

# Finding LibC

- The trick is to notice that we have the address of several libc functions in the GOT already
- For example, after we call puts, the linker puts the address of puts in libc into the GOT for us:

```
(gdb) x/i 0x400715
0x400715 <main+31>: call 0x400580 <puts@plt>
(gdb) x/i 0x400580
0x400580 <puts@plt>: jmp QWORD PTR [rip+0x200a92] # 0x601018
(gdb) x/xg 0x601018
0x601018: 0x00007ffff7a7c690
```



**This address is in libc!**

# Leaking a LibC address

- We can do the equivalent of `puts(puts@GOT)` to *leak* an address from libc
- Why does this help us?
- Once we know the address of `puts`, we can look up where `puts` is in libc
- Then just subtract off the offset of `puts`, and we have the base address of libc in memory
- From here we can get address of any other function in libc!

# Calculating libc Addresses

- In pwntools:

```
In [1]: from pwn import *
```

```
In [2]: # Suppose we know the address of puts is 0x00007ffff7a7c690
```

```
In [3]: puts_addr = 0x00007ffff7a7c690
```

```
In [4]: libc = ELF('libc-2.19.so')  
[*] '/home/moyix/offsec/week_9/rop/libc-2.19.so'  
Arch:      amd64-64-little  
RELRO:     Partial RELRO  
Stack:     Canary found  
NX:        NX enabled  
PIE:       PIE enabled
```

```
In [5]: libc_base = puts_addr - libc.symbols['puts']
```

```
In [6]: print hex(libc_base)  
0x7ffff7a0c930
```

```
In [7]: system_addr = libc_base + libc.symbols['system']
```

```
In [8]: print hex(system_addr)  
0x7ffff7a52ec0
```

# Bonus of having libc: More Gadgets!

## Main Binary

```
(pwn) moyix@lorenzo:~/offsec/week_9/rop$ ROPgadget --binary ./rop
Gadgets information
=====
0x0000000000400564 : adc byte ptr [rax], ah ; jmp rax
[...]
Unique gadgets found: 70
```

## libc

```
(pwn) moyix@lorenzo:~/offsec/week_9/rop$ ROPgadget --binary ./libc-2.19.so
Gadgets information
=====
0x000000000018a851 : adc ah, bh ; call qword ptr [rax - 0x1f0003ec]
[...]
Unique gadgets found: 25925
```

## libc also has lots of useful strings!


```
In [2]: libc = ELF('libc-2.19.so')
```

```
In [3]: libc.data.find('/bin/sh\x00')
```

```
Out[3]: 1574147
```

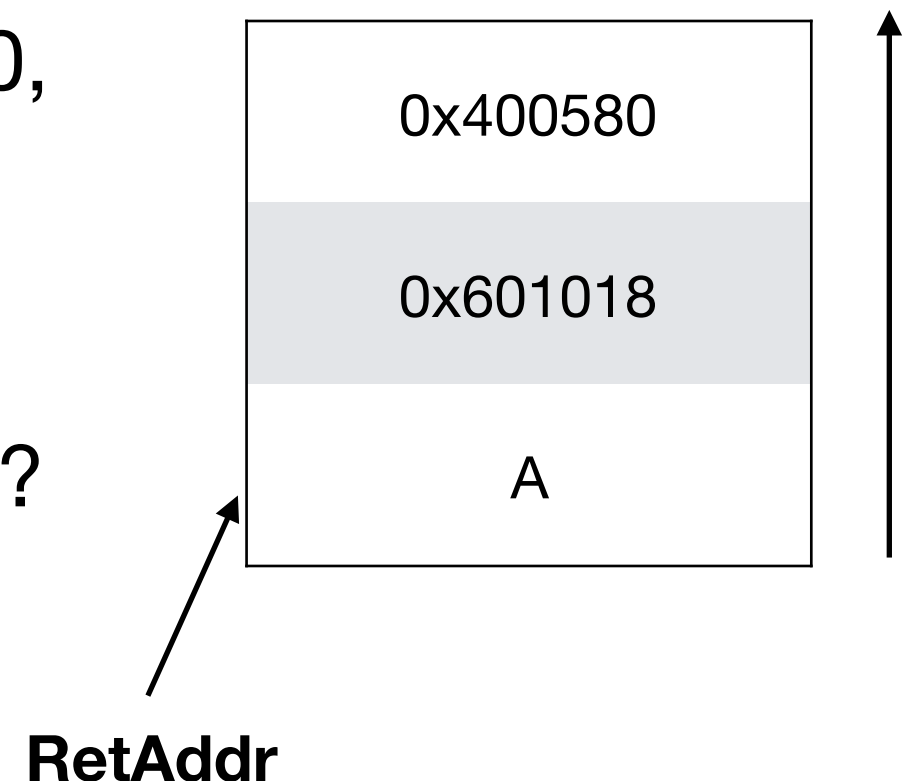
# Leaking puts

- Suppose we have a gadget  
    A  
pop rdi  
ret
- The PLT entry for puts is at 0x400580,  
and the GOT entry for puts is at  
0x601018
- What should our ROP chain look like?



# Leaking puts

- Suppose we have a gadget  
    A  
pop rdi  
ret
- The PLT entry for puts is at 0x400580,  
and the GOT entry for puts is at  
0x601018
- What should our ROP chain look like?



# Proposed Exploit

- Last time we proposed this exploit:
  - We have full control of the stack
  - The program did `gets(buf)`, or something similarly vulnerable
- We start our ROP...
  - `puts(puts_got_address)` to leak it => obtain address of `puts` in `libc`
  - We subtract the offset of `puts` in `libc` from that address to get the base address of `libc`
  - Then we add the offset of `system` in `libc` to get the real address of `system`
  - We jump there, with an argument of `"/bin/sh"`, and get a shell



# Problem!

- The exploit described on the last slide has a fatal flaw
- Do you see what it is?

# Problem!

- The exploit described on the last slide has a fatal flaw
- Do you see what it is?
  - After we return from `puts(puts_got_address)` the program crashes!
  - So although we can find out the address of `libc`, we can't do anything with it
  - Next time we run the program, the address of `libc` will be different – so our knowledge is useless!

# Solution: Return to main

- After our initial ROP chain that leaks the address of puts, we add the address of the main() function
- When the program finishes printing out the address, it will jump to the beginning of main... effectively restarting the program (but *without* changing where libc is)
- Now we can **re-exploit the program**
  - And we can make use of the libc addresses we learned the first time around!

# Advanced ROP

- You can't always get what you want
- Sometimes your binary won't have all the gadgets you need
- Example: need to do a 3 argument syscall, but we only have  
pop rdi; ret  
pop rsi; ret  
syscall; ret
- How can we get something into rdx?

# Some Ways to Set RDX

- Is there a `mov rdx, ____` and a way to set \_\_\_\_ ?
  - E.g. `mov rdx, rsi` since we already have a `pop rsi`
- What is `rdx` at the point where our ROP chain starts?
  - Can we influence it?
- Can we do arithmetic in the ROP chain?
  - `add`, `xor`, `inc`, etc gadgets
  - Last resort, but possible for "nice" numbers or numbers near `rdx`

# Getting Useful Strings and Constants

- We often need to get some useful string or constant into a known location (classic example: “/bin/sh\x00”)
  - May not have a libc leak
  - May not have a stack leak
- Is there some way we can put our own data into a place that has some writable memory at a fixed address?

# The .bss and .data Sections

- Every program will have a section of the executable dedicated to storing global variables
- Initialized data goes into .data, uninitialized data goes into .bss
- Recall: operating systems allocate memory in units of a *page* (4096 bytes on x86)
- This means there is often “slack” space at the end of one of these sections!

# .bss Example

```
.data (PROGBITS) section started {0x601050-0x601060}
00601050  __data_start:
00601050  00 00 00 00 00 00 00 00 00  .....
00601058  __dso_handle:
00601058                00 00 00 00 00 00 00 00  .....
.data (PROGBITS) section ended {0x601050-0x601060}
-----
.bss (NOBITS) section started {0x601060-0x601090}
00601060  stdout:
00601060  00 00 00 00 00 00 00 00 00-00 00 00 00 00 00 00 00  .....
00601070  stdin:
00601070  00 00 00 00 00 00 00 00 00  .....
00601078  completed.7585:
00601078                00 00 00 00 00 00 00 00  .....
00601080  buf:
00601080  00 00 00 00 00 00 00 00 00  .....

00601088  int64_t data_601088 = 0x0
.bss (NOBITS) section ended {0x601060-0x601090}
```

- Section “ends” at 0x601090
- ...but the rest of the page is mapped and writable
- So we have 0x601090-0x602000 for our own data



# Reading Data

- How do we actually put data there?
- ROP, of course!
- Assuming the program takes input, it must have a function imported like `fgets`, `gets`, or `read`
- So we can create a ROP chain that calls that function to read data into some space in the `.bss`
- Since the `.bss` is at a fixed location, this lets us put whatever data we want and know where it is

# Stack Pivots

- We saw one way (return to main) that you can make use of information you got from a ROP chain to continue your exploit
- Another way is a *stack pivot*:
  - We can call read() or similar to write a *second ROP chain* somewhere in the .bss
  - Then find a stack pivot gadget that allows us to **set the value of rsp** to point to our new ROP chain  
0x4006ad : **pop rsp** ; pop r13 ; pop r14 ; pop r15 ; ret
  - Everything after the **pop rsp** will happen on our **new stack**


# Stack Pivot Warning

- The .bss section is not that large
- And it's where your stack is now
- As a result, calling complicated functions that use a lot of stack space could make your program crash
  - Or start overwriting your ROP chain
- Best to keep things simple and avoid using the stack too much!

# One Last Trick

- If you have figured out where libc is but don't want to (or can't) construct a ROP chain to call `system("/bin/sh")`
- It turns out there are places within libc that if you jump to them will do the work of executing a shell for you!
  - In particular, partway through the `system()` function
- There are tools to find these locations, like `one_gadget`  
[https://github.com/david942j/one\\_gadget](https://github.com/david942j/one_gadget)

# one\_gadget

```
david942j at ~/one_gadget on master via  v2.5.3
```

```
→ one_gadget /lib/x86_64-linux-gnu/libc.so.6
```

```
0x4f2c5 execve("/bin/sh", rsp+0x40, environ)
```

```
constraints:
```

```
rcx == NULL
```

```
0x4f322 execve("/bin/sh", rsp+0x40, environ)
```


```
constraints:
```

```
[rsp+0x40] == NULL
```

```
0x10a38c execve("/bin/sh", rsp+0x70, environ)
```

```
constraints:
```

```
[rsp+0x70] == NULL
```

```
david942j at ~/one_gadget on master via  v2.5.3
```

```
→ 
```