

Heap Metadata Exploits

CS-UY 3943-G / CS-GY 9223-H

Last Time

- Looked at the data structures of the glibc allocator (ptmalloc2)
- Looked at:
 - Heap overflows
 - Use after free
- The exploits we discussed involved corrupting *application* data on the heap

This Time

- Introduction to heap metadata exploits
- Two case studies:
 - Exploiting unlink()
 - One-byte poison null overwrite

Metadata Exploits

- Many allocators use *inline metadata* – information about each heap chunk is stored alongside the application data
- The basic idea:
 - Corrupt metadata via heap overflow/type confusion
 - When the allocator interacts with it (e.g. during a malloc() or free()) it will have some controlled effect
 - Often *write-what-where*
- We can use this as an exploit primitive (GOT overwrite, function pointer overwrite, start a ROP chain, ...)

General Concept: Bins

- One thing to be aware of is that for efficiency, many allocators have different pools (***bins***) of free memory chunks
- These are usually divided up by size
- Different bins may use different data structures or allocation strategies
- Where your memory comes from when you do a malloc() depends on the size and allocations that have been made so far
- Good resource:
https://heap-exploitation.dhavalkapil.com/diving_into_glibc_heap/bins_chunks.html

Bin Types

- **Fast bins:** 10 singly-linked lists, one each for chunk sizes 16, 24, 32, 40, 48, 56, 64, 72, 80 and 88 bytes
- **Small bins:** 62 doubly-linked lists, again one for each size, where the sizes are 16, 24, ... , 504 bytes
- **Large bins:** 63 doubly-linked lists, one for each size in the ranges [512 - 568], [576 - 632], ...
- **Unsorted bin:** a single bin that stores recently-freed small and large bins. Chunks may end up here temporarily in an attempt to reuse recently-freed allocations quickly.

More Terminology

- **Top chunk:** the chunk at the top of the arena (i.e., the top of the currently allocated heap area for this thread)
 - If there's no other space, the allocator will try to use this chunk
 - If there's still no space, it will grow the heap by extending the top chunk
- **Last remainder chunk:** if a chunk of the exact size requested doesn't exist, a chunk may be split in two. The unused half will be set as the "last remainder chunk"

unlink()

- As a warm-up, we'll look at exploiting the way the unlink() function works
- This is an older technique that no longer works due to extra sanity checks added in the allocator
- But the mechanisms are simple so it's still good as a toy example

Heap Chunk Structure

- For a free chunk, we have:

```
struct malloc_chunk {  
    INTERNAL_SIZE_T  
    INTERNAL_SIZE_T  
    struct malloc_chunk*  
    struct malloc_chunk*  
};
```

mchunk_prev_size;

mchunk_size;

fd;

bk;

Size of previous chunk

Size of this chunk

Pointer to previous free chunk

Pointer to next free chunk

Heap Chunk Structure

- For a free chunk, we have:

```
struct malloc_chunk {  
    INTERNAL_SIZE_T  
    INTERNAL_SIZE_T  
    struct malloc_chunk*  
    struct malloc_chunk*  
};
```

```
mchunk_prev_size;  
mchunk_size;  
fd;  
bk;
```

Note: *prev_size* also stores a flag that says whether the previous chunk is actually free!

Size of previous chunk


Size of this chunk

Pointer to previous free chunk

Pointer to next free chunk

unlink() Macro


```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```



Intended Use

unlink() Macro


```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```



Intended Use

unlink() Macro


```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```



Intended Use

unlink() Macro

```
#define unlink( P, BK, FD ) {  
    BK = P->bk;  
    FD = P->fd;  
    FD->bk = BK;  
    BK->fd = FD;  
}
```



Intended Use

unlink() operations

- Suppose control the content of the heap chunk **P**
- What do the unlink operations do?
- Suppose:
 - chunk.BK is at offset 12 in the struct
 - chunk.FD is at offset 8 in the struct

unlink() operations

- Then:
 - $FD \rightarrow bk = BK$ is the same as $*(FD+12) = BK$
 - “Write BK to the address $FD + 12$ ”
 - $BK \rightarrow fd = FD$ is the same as $*(BK+8) = FD$
 - “Write FD to the address $BK + 8$ ”
- Since we control both BK and FD, this is a **write-what-where**

Exploit Example

- Suppose we want to overwrite a GOT pointer at 0x602020 (puts) with the address 0x4006e0
- We set $FD = 0x602020 - 12 = 0x602014$
- We set $BK = 0x4006e0$
- Then when we `unlink()`, $FD \rightarrow bk = BK$ will do $*(FD + 12) = *(0x602020) = 0x4006e0$

Exploit Example

- Suppose we want to overwrite a GOT pointer at 0x602020 (puts) with the address 0x4006e0
- We set $FD = 0x602020 - 12 = 0x602014$
- We set $BK = 0x4006e0$
- Then when we `unlink()`, $FD \rightarrow bk = BK$ will do $*(FD + 12) = *(0x602020) = 0x4006e0$

Note: this doesn't *quite* work because right afterward $BK \rightarrow fd = FD$ will try to write to a read-only code page!

Getting to unlink()

- Now that we understand the basic primitive, we still need to get the allocator to call **unlink()** on data we control
- Where is unlink() called?
 - Inside malloc(), to grab a chunk from the free list and use it for an allocation
 - Inside free(), to consolidate adjacent free chunks

Getting to unlink()

- Now that we understand the basic primitive, we still need to get the allocator to call **unlink()** on data we control
- Where is unlink() called?
 - Inside malloc(), to grab a chunk from the free list and use it for an allocation
 - **Inside free(), to consolidate adjacent free chunks**

Chunk Consolidation

- The code to consolidate a chunk with the previous one looks like:

```
/* consolidate backward */
if (!prev_inuse(p)) {
    prevsize = prev_size (p);
    size += prevsize;
    p = chunk_at_offset(p, -((long) prevsize));
    unlink(p, bck, fwd);
}
```

Chunk Consolidation

- The code to consolidate a chunk with the previous one looks like:

```
/* consolidate backward */  
if (!prev_inuse(p)) {  
    prevsize = prev_size (p);  
    size += prevsize;  
    p = chunk_at_offset(p, -((long) prevsize));  
    unlink(p, bck, fwd);  
}
```

Checks "in use" bit

Gets address of previous chunk

unlink()

Exploit Scenario

- Suppose we can set up the heap to look like:



- And there is an overflow that lets us write past the end of Chunk 1, letting us overwrite Metadata and Chunk 2:



- Then if the program tries to free() Chunk 2, and we can create **fake** metadata that makes the allocator think Chunk 1 is also free, we can get it to call unlink() on data we control

Fake Chunk

- Interpreted as a free chunk, the fields we control look like



- We want to set them to values that will result in the allocator calling `unlink()` on our fake chunk
- We can do this by setting “in-use” (U) to 0, `prev_size` to 0, and `size` large enough that it will not go in the fast bins



Consolidation

- When free() tries to consolidate the chunk, let's look at what the code sees now:

```
/* consolidate backward */  
if (!prev_inuse(p)) {  
    prevsize = prev_size (p);  
    size += prevsize;  
    p = chunk_at_offset(p, -((long) prevsize));  
    unlink(p, bck, fwd);  
}
```

“In use” is 0 so we will try to consolidate

Since prev_size is 0, the address of the “previous” chunk will end up being the start of our fake chunk

**Allocator will call unlink() on our fake chunk
Assuming we set FD and BK appropriately, this will
execute our write-what-where!**

Mitigations

- The unlink() trick described here no longer works
- Two checks were added:
 - Checking that the next chunk's prev_size equals this chunk's size
 - Checking the linked list pointers:
 - FD->bk points to current chunk
 - BK->fd points to current chunk

Off-by-One Overwrites

- In many programs you will find “off-by-one” vulnerabilities
 - Usually from calculating the size of a buffer wrong, e.g. by forgetting about the byte for the terminating null
- This may seem fairly useless – what could you do with just a single byte overwrite?
- Even more restriction: very often the overwrite will only let you write a “0” (for the null terminator)

Poison NULL

- Back in 1998, Olaf Kirch showed that even one-byte NULL overwrites could be dangerous
 - Demonstrated a one-byte NULL overwrite that modified the saved base pointer on the stack, eventually leading to full code execution
- In 2014, Tavis Ormandy of Google Project Zero showed that a one-byte NULL overwrite could *also* be used for a heap metadata exploit in glibc

`glibc` Poison NULL


- The basic trick is that a one-byte overflow that writes a 0 will have the effect of setting the least significant byte of the next chunk's size field to 0
- This will effectively *shrink* the size of that chunk from the allocator's point of view
 - Any calculations done by the allocator with this incorrect size will be affected!
- We're going to use this to eventually create two *overlapping chunks*

Shrinking a Chunk


- Allocate three consecutive chunks: A, B, C
- Free B
- Overflow from A, making B.size smaller
- Allocate two new chunks, B1 and B2 in the free space
 - **C's prev_size is not updated correctly (why?)**

Shrinking a Chunk


- Free B1 and C
 - When we free C, the allocator thinks B is still free because C's prev_size was not updated!
 - **The allocator will merge C and B to create one large free area**
- Now we allocate a final object. It will be placed at the start of **B** – *overlapping with **B2***



Initial state



B is free



Overflow: size(B) = 0x200

Overflow into B


- Size truncated to 0x200 from 0x208
- Further allocations in that space do not properly update C's "prev_size" field




Two allocations within the old B chunk
The first is not a fastbin



The beginning of the old B chunk is free



C is freed and merged with the old B, where
a valid non-fastbin free chunk resides




1+ allocations larger than B1's initial size
B2 is overlapped

Source: https://www.contextis.com/media/downloads/Glibc_Adventures_The_forbidden_chunks.pdf

Using the Overlap

- What can we do with two overlapping heap chunks?
- If we can make some (allowed) modification to a field of one of the overlapping objects, we can change a *different* field in the other



- By editing Obj1.buffer, we can set Obj2.pointer to whatever we want!

Lots More Metadata Attacks

- There are many more types of heap metadata exploit techniques, even just with glibc
 - House of _____ (force, spirit, Einherjar, orange, lore, ...)
 - Unsorted bin attacks
 - Exploits using the “tcache” feature of glibc
- Details: <https://github.com/shellphish/how2heap>